

Monte Carlo Tree Search

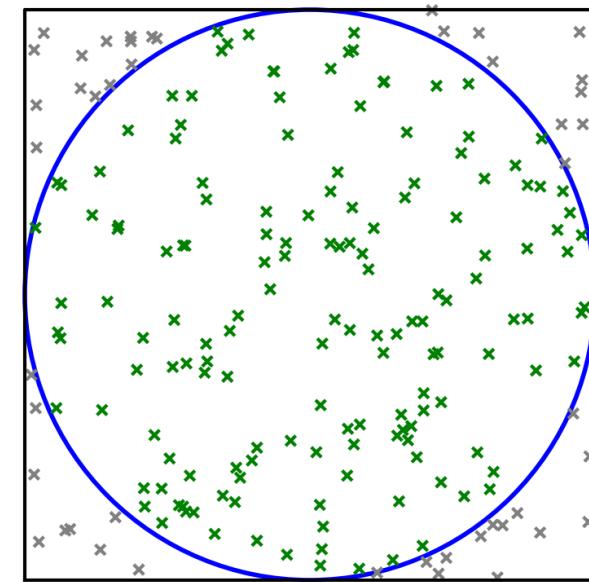
Monte Carlo

- ▶ It refers to algorithms using random numbers (pseudo random variables) to solve computation problems.
- ▶ Is it the name of a person? No, it is the name of a city famous for gambling.
- ▶ Classic applications:
 - ▶ Estimation of π
 - ▶ Monte Carlo integration

Monte Carlo Application: Estimation Of π

- ▶ Generate random points (x, y) in $[0,1]$.
- ▶ Count how many points fall inside the circle centered at $(0.5,0.5)$ with radius 0.5.
- ▶ $\pi \approx 4 \times (\text{points inside circle} / \text{total points})$.

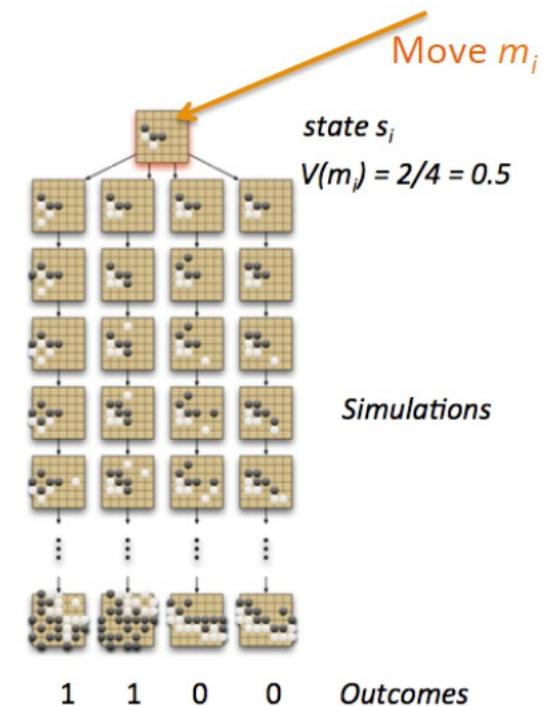
Monte Carlo: Estimating π
 $\pi \approx 4 * (\text{points inside circle} / \text{total points})$



Trying to Apply Monte Carlo to Games

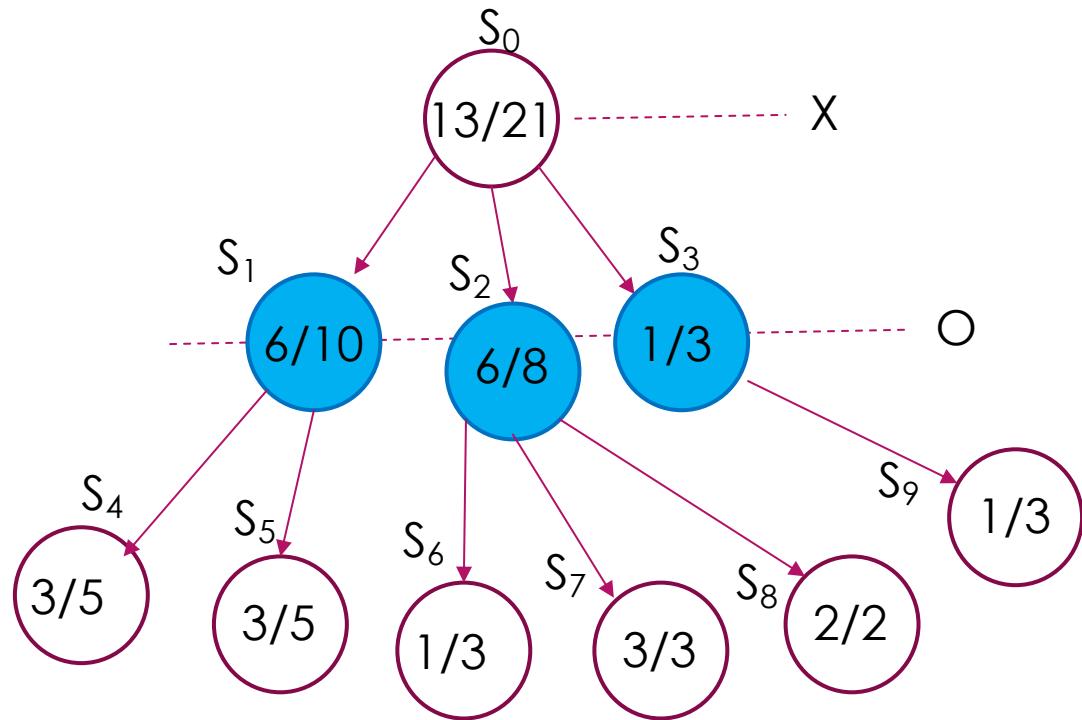
Naive Idea:

- ▶ Run many random play-outs of the game.
- ▶ Estimate the win rate of each move by averaging outcomes.
- ▶ Choose the move with the highest estimated win rate.



Limits

- ▶ Random play-outs ignore the decision structure of the game tree.
Example: X will select node $S_1(6/8)$, but as X's opponent O, O may select $S_6(1/3)$, not $S_7(3/3)$ or $S_8(2/2)$.
- ▶ Some moves are explored too little, others too much.
Example: $S_3(1/3)$ only explored 3 times while its brother S_1 explored 10, S_2 explored 8.
- ▶ No mechanism to balance exploration vs exploitation.



How to improve

- ▶ Incorporating adversarial thinking, assumes an opponent who minimizes our advantage, and thus searches for robust strategies.
- ▶ Prioritizes promising regions for precise evaluation, rapidly dismisses unpromising areas with minimal simulation.
- ▶ Balance between exploration and exploitation.

MCTS can do all above!

Monte Carlo Tree Search (MCTS)

◆ A Game AI Algorithm

- Decision Making under Uncertainty
- widely used in AlphaGo, Game Bots

◆ Core Idea

- Evaluate moves through random simulations
- Build search tree incrementally
- Balance exploration vs exploitation

◆ Advantages

- No domain knowledge required
- Easy parallelization
- Anytime algorithm (can stop anytime)

MCTS Basic Process

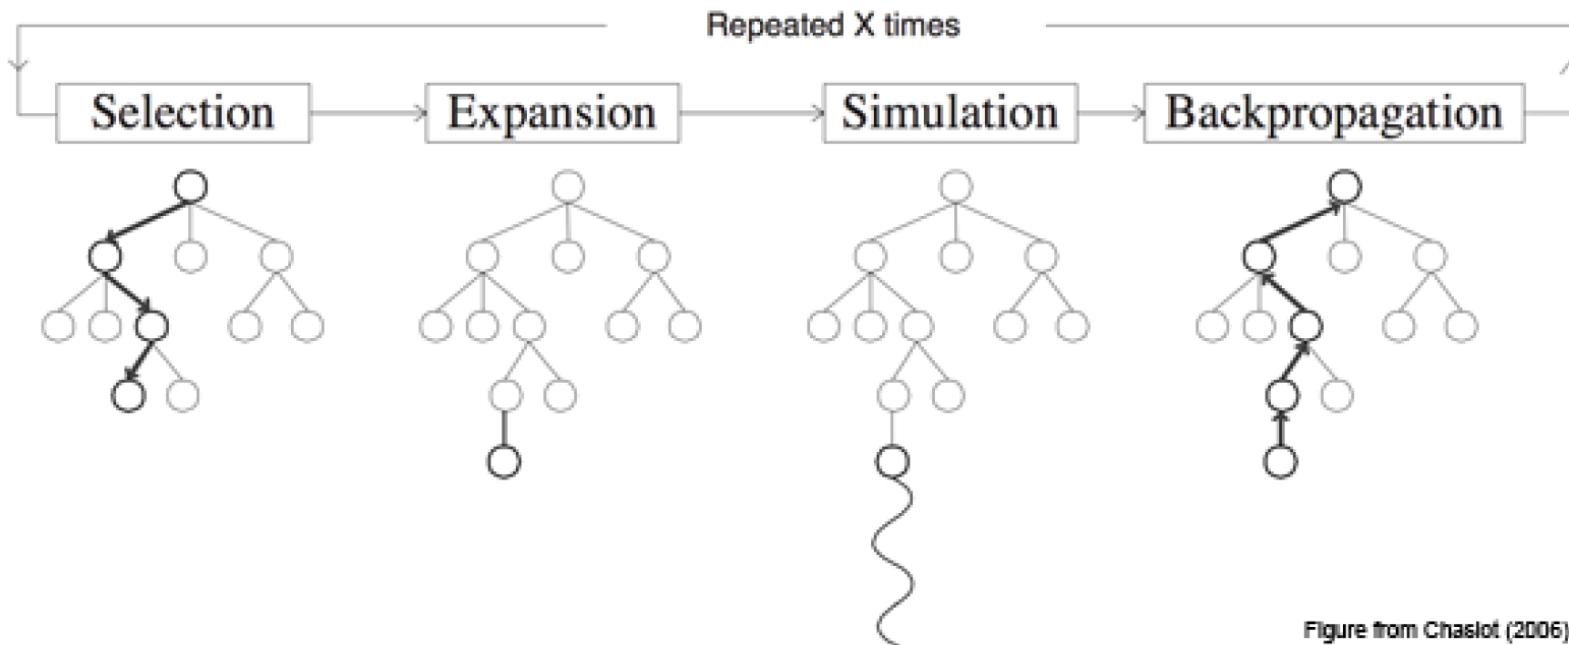


Figure from Chaslot (2006)

- ▶ Selection: from root node, recursively apply a **child selection policy** to descend through the tree until find the most urgent expandable leaf node.
- ▶ Expansion: Add one or more child nodes.
- ▶ Simulation: Run random simulation from the new node to terminal state.
- ▶ Backpropagation: The simulation result is "backed up" through the selected nodes to update their statistics.

Example: Selection and Expansion

Initialization:



Selection:

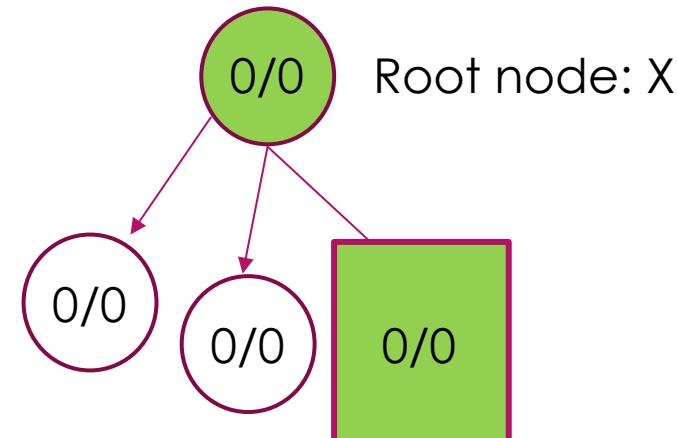


Current node has no child info, select itself.

What selection policy to select a child if it has child?

Expansion:

Add one or more child nodes.
Return a new node.



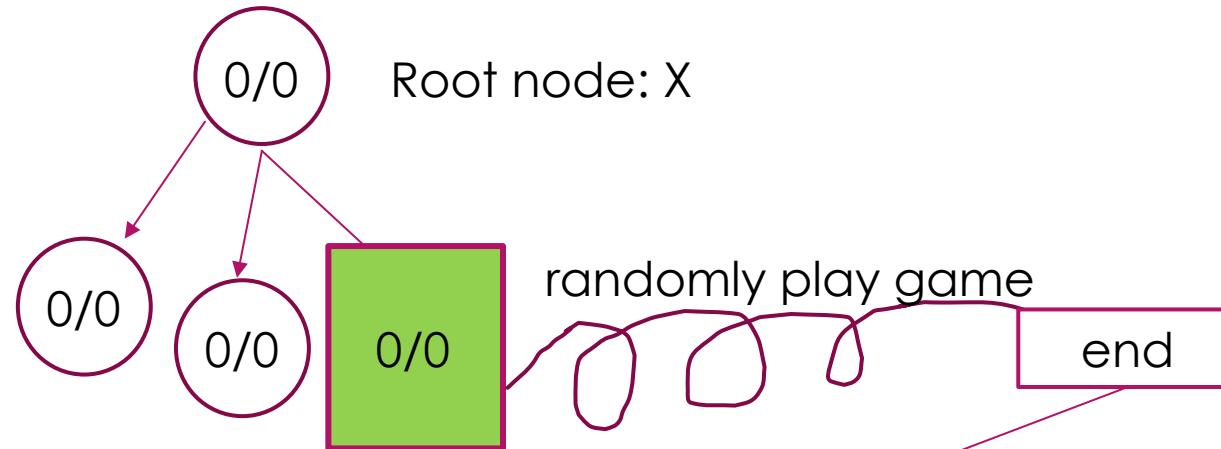
If there's more than one child, pick one at random or by selection policy.

Example: Simulation and Backpropagation

Simulation

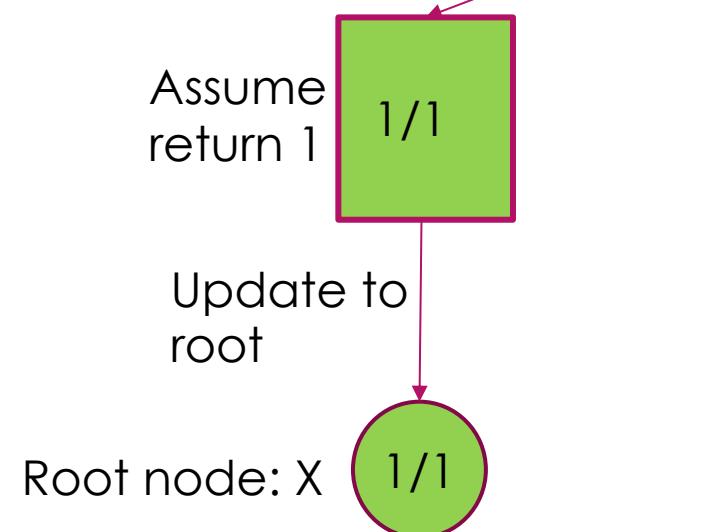
Use **random strategy** to play from the child state until game ends.

Return the result of the simulated game.

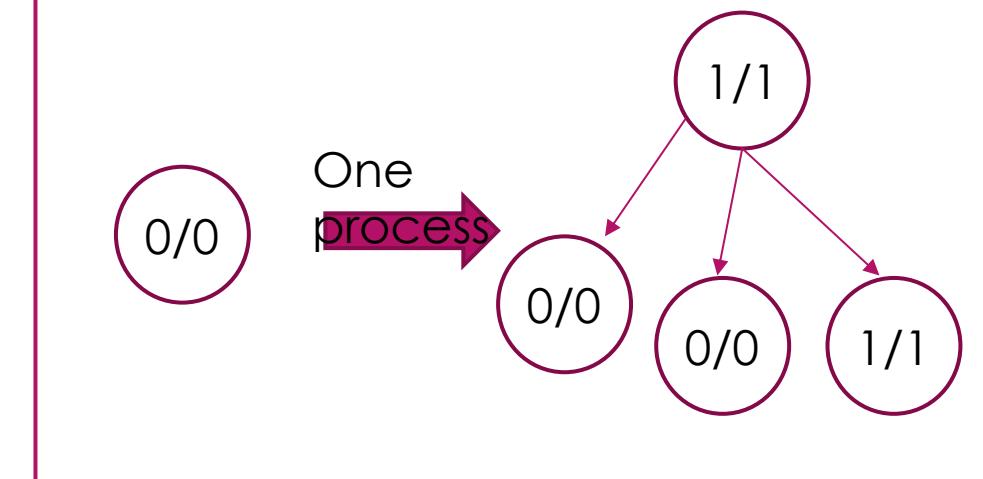


Backpropagation

Update win rates and visit counts **from leaf to root**.

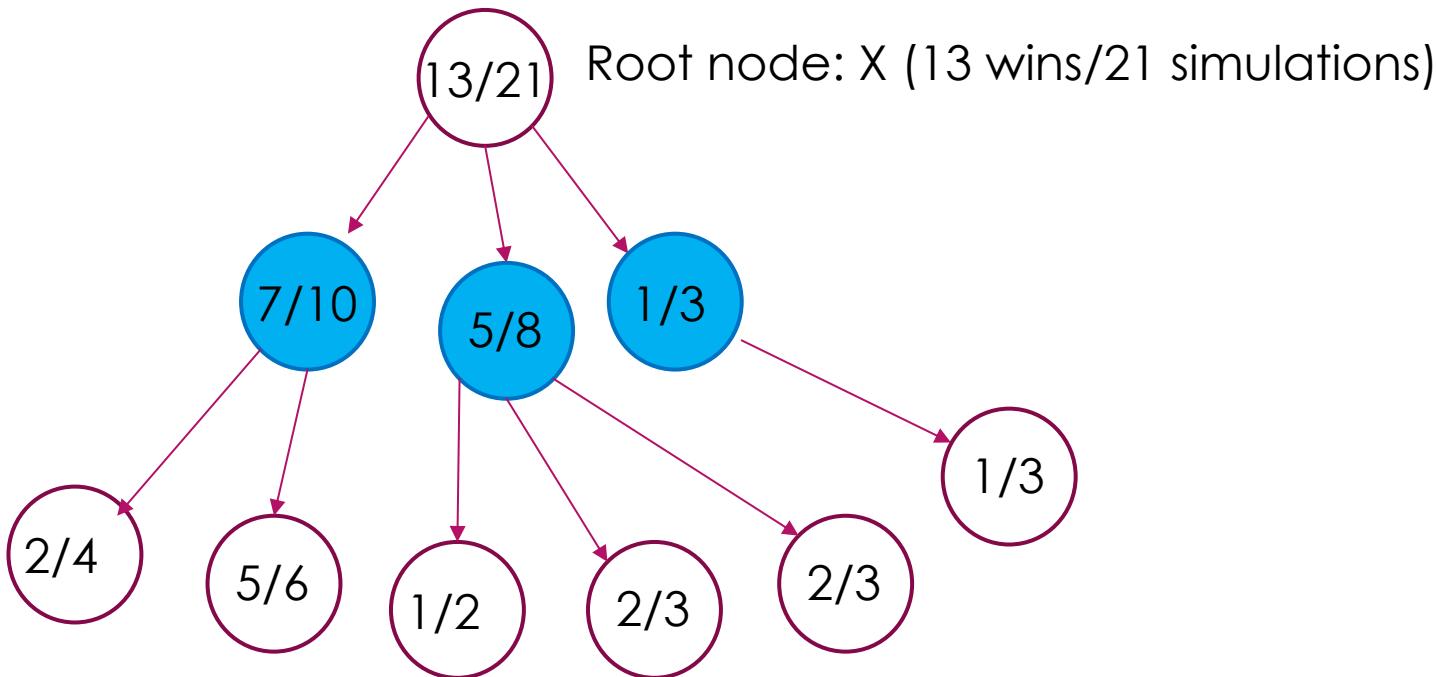


Return 1 for X win, 0 for X loss



Simulation Results Statistics

Assume the process repeat **21** times, and X win **13** times:



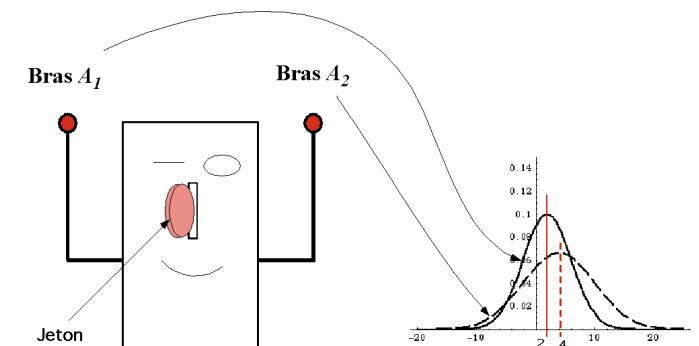
Based on the above tree, we go on to discuss the **selection policy**.

Multi-Armed Bandit Problem(MAB)



Assumption

- ▶ Choice among several arms
- ▶ Each arm pull is independent of other pulls
- ▶ Each arm has unknown reward distribution



Which arm has the best average payoff?

Regret in Multi-Armed Bandit (MAB)

- ▶ Optimal strategy: always pull the best arm
- ▶ Actual strategy: mix of exploration and exploitation
- ▶ **Regret = Optimal reward – Actual reward**

Example: Regret

A



$P(A \text{ wins}) =$
60%

B



$P(B \text{ wins}) =$
55%

C



$P(C \text{ wins}) =$
40%

A sequence of pulls:

Pull A 300 times, win 170 times

Pull B 300 times, win 150 times

Pull C 400 times, win 165 times

Best sequence of pulls:

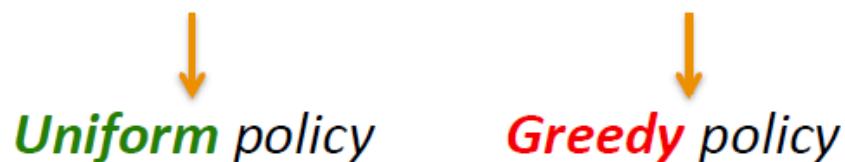
Pull A 1000 times, win 600 times

Regret = $60\% - \frac{485}{1000} =$
0.115

How To Minimize regret ?

- ▶ Actual strategy: mix of exploration and exploitation
 - ▶ Exploitation: choose arm with the highest win rate
 - ▶ Exploration: explore the arm with the highest uncertainty

Need to **balance** exploration and exploitation



Upper Confidence Bound (UCB)

- ▶ UCB1 formula[Aueretal.2002]
 - ▶ First,try each arm once
 - ▶ Then at each time step, choose arm i that maximizes the UCB1 formula:

$$v_i + C \times \sqrt{\frac{\ln(N)}{n_i}}$$

Annotations for the UCB1 formula:

- v_i (value estimate)
- C (tunable parameter)
- $\ln(N)$ (total number of trials)
- n_i (num trials for arm i)

Very sensitive to C

Upper Confidence Tree(UCT)=MCTS+UCB

Apply UCB to MCTS selection phase:

$$\frac{U_i}{N_i} + C \sqrt{\frac{\ln(N)}{N_i}}$$

U_i : Number of wins for the child node i

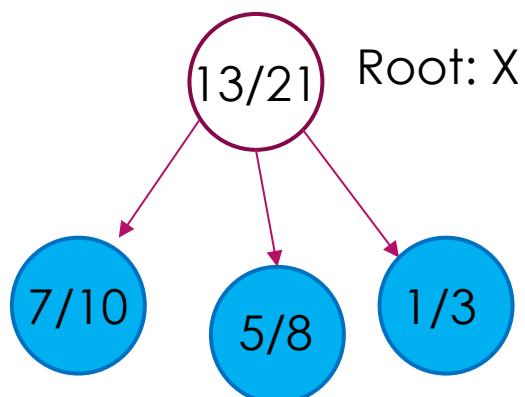
N_i : Number of simulations for the child node i

N : Total number of simulations

C : coefficient (tunable parameter), exploration-exploitation trade-off

UCT Calculation Example (X Turn)

Root node: $N = 21$



set $C = 10$
child node 1:
 $UCT = 7/10 + 10 * \sqrt{\ln(21)/10}$
 ≈ 6.22

child node 2:
 $UCT = 5/8 + 10 * \sqrt{\ln(21)/8}$
 ≈ 6.80

child node 3:
 $UCT = 1/3 + 10 * \sqrt{\ln(21)/3}$
 ≈ 10.40 (selected, high C favors exploration)

set $C = 0.5$
child node 1:
 $UCT = 7/10 + 0.5 * \sqrt{\log(21)/10}$
 ≈ 0.976 (selected, low C favors exploitation)

child node 2:
 $UCT = 5/8 + 0.5 * \sqrt{\log(21)/8}$
 ≈ 0.934

child node 3:
 $UCT = 1/3 + 0.5 * \sqrt{\log(21)/3}$
 ≈ 0.836

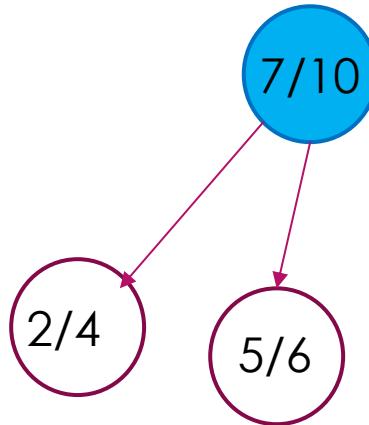
Assume set C = 10, from the root,
select the child :

7/10

Now, current parent node: N = 10

child node 1: $UCB = 2/4 + 10 * \sqrt{\ln(10)/4} \approx 8.09$

child node 2: $UCB = 5/6 + 10 * \sqrt{\ln(10)/6} \approx 7.03$



Right? **No**

UCT Calculation Example (Opponent's Turn)

During opponent's turn, invert win rate in UCB formula:

$$(1 - \frac{U_i}{N_i}) + C \sqrt{\frac{\ln(N)}{N_i}}$$

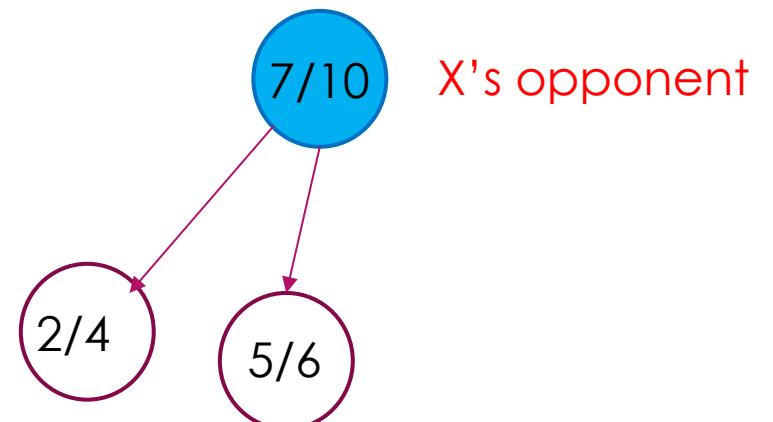
Because opponent wants to minimize your win rate.

Calculation Example:

Current parent node: $N = 10$

child node 1: $UCB = (1 - 2/4) + 10 * \sqrt{\ln(10)/4} = 8.09$

child node 2: $UCB = (1 - 5/6) + 10 * \sqrt{\ln(10)/6} = 6.36$



Limitations of MCTS

- ▶ Quality depends on simulations
- ▶ May need many iterations
- ▶ Parameter sensitivity
- ▶ Memory intensive

Summary

- ✓ MCTS is a decision-making method requiring no prior knowledge.
- ✓ Applications: Game AI, path planning, automated reasoning
- ✓ UCT balances exploration and exploitation for efficient decision-making
- ✓ Future Directions: Deep learning integration, Real-time applications, Multi-agent systems.